

---

# Machine Learning

## Topic: Linear Discriminants

Bryan Pardo, EECS 349 Machine Learning, 2021

Thanks to Mark Cartwright for his contributions to these slides  
Thanks to Alpaydin, Bishop, and Duda/Hart/Stork for images and ideas

# Recall: Regression Learning Task

---

There is a set of possible examples  $X = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$

Each example is a **vector** of  $k$  **real valued attributes**

$$\mathbf{x}_i = \langle x_{i1}, \dots, x_{ik} \rangle$$

A target function maps  $X$  onto some **real value**  $Y$

$$f : X \rightarrow Y$$

The DATA is a set of tuples  $\langle \text{example}, \text{response value} \rangle$

$$\{\langle \mathbf{x}_1, y_1 \rangle, \dots, \langle \mathbf{x}_n, y_n \rangle\}$$

Find a **hypothesis**  $h$  such that...

$$\forall \mathbf{x}, h(\mathbf{x}) \approx f(\mathbf{x})$$

# Discrimination Learning Task

---

There is a set of possible examples  $X = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$

Each example is a **vector** of  $k$  **real valued attributes**

$$\mathbf{x}_i = \langle x_{i1}, \dots, x_{ik} \rangle$$

A target function maps  $X$  onto some **categorical variable**  $Y$

$$f : X \rightarrow Y$$

The DATA is a set of tuples  $\langle \text{example}, \text{response value} \rangle$

$$\{\langle \mathbf{x}_1, y_1 \rangle, \dots, \langle \mathbf{x}_n, y_n \rangle\}$$

Find a **hypothesis**  $h$  such that...

$$\forall \mathbf{x}, h(\mathbf{x}) \approx f(\mathbf{x})$$

# Reminder about notation

---

- $\mathbf{x}$  is a vector of attributes  $\langle x_1, x_2, \dots, x_k \rangle$
- $\mathbf{w}$  is a vector of weights  $\langle w_1, w_2, \dots, w_k \rangle$
- Given this...

$$g(x) = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_k x_k$$

- We can notate it with linear algebra as

$$g(x) = w_0 + \mathbf{w}^T \mathbf{x}$$

# Recall: $w_0$

---

- $g(x) = w_0 + \mathbf{w}^T \mathbf{x}$  is ALMOST what we want, but that pesky offset  $w_0$  is not in the linear algebra part yet.
- If we define  $\mathbf{w}$  to include  $w_0$  and  $\mathbf{x}$  to include an  $x_0$  that is always 1, now...

$\mathbf{x}$  is a vector of attributes  $\langle 1, x_1, x_2, \dots, x_k \rangle$


$\mathbf{w}$  is a vector of weights  $\langle w_0, w_1, w_2, \dots, w_k \rangle$

- This lets us notate things as...

$$g(x) = \mathbf{w}^T \mathbf{x}$$


# Visually: Where to draw the line?

---



# Two-Class Classification

$g(\mathbf{x}) = 0$  defines a decision boundary that splits the space in two






If a line exists that does this without error, the classes are *linearly separable*

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } g(\mathbf{x}) > 0 \\ -1 & \text{otherwise} \end{cases}$$

# Example 2-D decision boundaries

---

$$0 = g(x) = w_0 + w_1 x_1 + w_2 x_2 = \mathbf{w}^T \mathbf{x}$$







# What's the difference?

---

$$0 = g(\mathbf{x}) = w_0 + w_1 x_1 + w_2 x_2 = \mathbf{w}^T \mathbf{x}$$

What's the difference between these two?



$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } g(\mathbf{x}) > 0 \\ -1 & \text{otherwise} \end{cases}$$

# Loss/Objective function

---

- To train a model (e.g. learn the weights of a useful line) we define a measure of the “goodness” of that model. (e.g. the number of misclassified points).
- We make that measure a function of the parameters of the model (and the data).
- This is called a loss function, or an objective function.
- We want to minimize the loss (or maximize the objective) by picking good model parameters.

# Classification via regression

---

- Linear regression's loss function is the the squared distance from a data point to the line, summed over all data points.
- The line that minimizes this function can be calculated by applying a simple formula.

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$


- Can we find a decision boundary in one step, by just repurposing the math we used for finding a regression line?

# Classification via regression


---

- Label each class by a number
- Call that number the response variable
- Derive closed-form regression solution
- Round the regression prediction to the nearest label number

# An example



# What happens now?



# Classification via regression take-away

---

- Closed form solution: simple formula for getting the regression line
- Residual sum of squares is a bad fit for classification: very sensitive to outliers
- What's the natural mapping from categories to the real numbers?

# What can we do instead?


---

- Let's define an objective (aka "loss") function that directly measures the thing we want to get right
- Then let's try and find the line that minimizes the loss.
- How about basing our loss function on the number of misclassifications?



# sum of squared errors (SSE)

---




$$g(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 = 0$$
$$= \mathbf{w}^T \mathbf{x}$$

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } g(\mathbf{x}) > 0 \\ -1 & \text{otherwise} \end{cases}$$

$$SSE = \sum_i^n (y_i - h(\mathbf{x}_i))^2$$

# sum of squared errors (SSE)

---



$$g(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 = 0$$
$$= \mathbf{w}^T \mathbf{x}$$

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } g(\mathbf{x}) > 0 \\ -1 & \text{otherwise} \end{cases}$$

$$SSE = \sum_i^n (y_i - h(\mathbf{x}_i))^2$$

# No closed form solution!

---

- For many objective functions we can't find a formula to get the best model parameters, like we could with regression.
- The objective function from the previous slide is one of those "no closed form solution" functions.
- This means we have to try various guesses for what the weights should be and try them out.
- Let's look at the perceptron approach.

# Let's learn a decision boundary

---

- We'll do 2-class classification
- We'll learn a linear decision boundary

$$0 = g(x) = \mathbf{w}^T \mathbf{x}$$

- Things on each side of 0 get their class labels according to the sign of what  $g(x)$  outputs.

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } g(\mathbf{x}) > 0 \\ -1 & \text{otherwise} \end{cases}$$


- We will use the Perceptron algorithm.

# The Perceptron



---

- Rosenblatt, F. (1958). The perceptron: A probabilistic model for information storage and organization in the brain. *Psychological Review*, 65(6), 386-408
- The “first wave” in neural networks
- A linear classifier


# A single perceptron




# Weights define a hyperplane in the input space



# Classifies any (linearly separable) data




**AND**







# Different logical functions are possible




OR



# And, Or, Not are easy to define




**NOT**



# One perceptron: Only linear decisions


---



This is XOR.

It can't learn  
XOR.

# Combining perceptrons can make any Boolean function



...if you can set the weights & connections right

# Defining our goal


---

$D$  is our data, consisting of training examples  $\langle \mathbf{x}, y \rangle$ . Remember  $y$  is the true label (drawn from  $\{1, -1\}$  and  $\mathbf{x}$  is the thing being labeled.

Our goal : make  $(\mathbf{w}^T \mathbf{x})y > 0$  for all  $\langle \mathbf{x}, y \rangle \in D$

Think about why this is the goal.

# An example



Goal: classify ● as +1 and ▲ as -1 by putting a line between them.

Our objective function is...

$$(\mathbf{w}^T \mathbf{x})y > 0$$


Start with a randomly placed line.

$$\mathbf{w} = [w_0, w_1, w_2] = [-5, 0, 1]$$

Measure the objective for each point.

Move the line if the objective isn't met.

# An example



Goal: classify ● as +1 and ▲ as -1 by putting a line between them.

Our objective function is...

$$(\mathbf{w}^T \mathbf{x})y > 0$$


Start with a randomly placed line.

$$\mathbf{w} = [w_0, w_1, w_2] = [-5, 0, 1]$$

●  $(\mathbf{w}^T \mathbf{x})y = [-5, 0, 1]^T [1, 5, 7](1)$   
 $= 2$   
 $> 0$

Objective met. Don't move the line.

# An example



Goal: classify ● as +1 and ▲ as -1 by putting a line between them.

Our objective function is...

$$(\mathbf{w}^T \mathbf{x})y > 0$$

Start with a randomly placed line.


$$\mathbf{w} = [w_0, w_1, w_2] = [-5, 0, 1]$$

$$\begin{aligned} \text{▲ } (\mathbf{w}^T \mathbf{x})y &= [-5, 0, 1]^T [1, 2, 6](-1) \\ &= (-5 + 6)(-1) \\ &= -1 \\ &< 0 \end{aligned}$$

Objective not met. Move the line.



# An example



Goal: classify  as +1 and  as -1 by putting a line between them.

Our objective function is...

$$(\mathbf{w}^T \mathbf{x})y > 0$$

Start with a randomly placed line.

$$\mathbf{w} = [w_0, w_1, w_2] = [-5, 0, 1]$$

Let's update the line by doing  $\mathbf{w} = \mathbf{w} + \mathbf{x}(y)$ .

$$\begin{aligned}\mathbf{w} = \mathbf{w} + \mathbf{x}(y) &= [-5, 0, 1] + [1, 2, 6](-1) \\ &= [-6, -2, -5]\end{aligned}$$

# Now what ?

---

- What does the decision boundary look like when  $\mathbf{w} = [-6, -2, -5]$  ? Does it misclassify the blue dot now?
- What if we update it the same way, each time we find a misclassified point?
- Could this approach be used to find a good separation line for a lot of data?

# Perceptron Algorithm

The decision boundary

$$0 = g(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

The classification function

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } g(\mathbf{x}) > 0 \\ -1 & \text{otherwise} \end{cases}$$

$m = |D| = \text{size of data set}$

The weight update algorithm

$\mathbf{w} = \text{some random setting}$

Do

$$k = (k + 1) \bmod(m)$$

$$\text{if } h(\mathbf{x}_k) \neq y_k$$

$$\mathbf{w} = \mathbf{w} + \mathbf{x}_k y_k$$

Until  $\forall k, h(\mathbf{x}_k) = y_k$


**Warning: Only guaranteed to terminate if classes are linearly separable!**

**This means you have to add another exit condition for when you've gone through the data too many times and suspect you'll never terminate.**

# Perceptron Algorithm

---

- Example:




Red is the positive class

Blue is the negative class

# Perceptron Algorithm

---

- Example (cont'd):




Red is the positive class

Blue is the negative class

# Perceptron Algorithm

- Example (cont'd):




Red is the positive class

Blue is the negative class

# Perceptron Algorithm

---


- Example (cont'd):



Red is the positive class

Blue is the negative class

# Multi-class Classification



When there are  $N$  classes you can classify using  $N$  discriminant functions.

Choose the class  $c$  from the set of all classes  $C$  whose function  $g_c(\mathbf{x})$  has the maximum output

Geometrically divides feature space into  $N$  **convex** decision regions

$$h(\mathbf{x}) = \operatorname{argmax}_{c \in C} g_c(\mathbf{x})$$



# Multi-class Classification

---

A class label

$$c = h(\mathbf{x}) = \underset{c \in \mathcal{C}}{\operatorname{argmax}} g_c(\mathbf{x})$$


Set of all classes

Remember  $g_c(\mathbf{x})$  is the inner product of the feature vector for the example  $(\mathbf{x})$  with the weights of the decision boundary hyperplane for class  $c$ . If  $g_c(\mathbf{x})$  is getting more positive, that means  $(\mathbf{x})$  is deeper inside its “yes” region.

Therefore, if you train a bunch of 2-way classifiers (one for each class) and pick the output of the classifier that says the example is deepest in its region, you have a multi-class classifier.

# Pairwise Multi-class Classification

If they are not linearly separable (singly connected convex regions), may still be pair-wise separable, using  $N(N-1)/2$  linear discriminants.



$$g_{ij}(\vec{x} | \vec{w}_{ij}, w_{ij0}) = w_{ij0} + \sum_{l=1}^K w_{ijl} x_l$$

choose  $C_i$  if  
 $\forall j \neq i, g_{ij}(\mathbf{x}) > 0$

---

# Appendix

(stuff I didn't have time to discuss in class...and for which I haven't updated the notation. )

# Linear Discriminants

---


- A linear combination of the attributes.

$$g(\vec{x} \mid \vec{w}, w_0) = w_0 + \vec{w}^T \vec{x} = w_0 + \sum_{i=1}^k w_i a_i$$

- Easily interpretable
- Are optimal when classes are Gaussian and share a covariance matrix

# Fisher Linear Discriminant Criteria

- Can think of  $\vec{w}^T \vec{x}$  as dimensionality reduction from K-dimensions to 1
- Objective:
  - Maximize the difference between class means
  - Minimize the variance within the classes



$$J(\vec{w}) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2}$$

where  $s_i$  and  $m_i$  are the sample variance and mean for class  $i$  in the projected dimension. We want to maximize  $J$ .

# Fisher Linear Discriminant Criteria

---

- Solution:


$$\vec{w} = \mathbf{S}_W^{-1}(\vec{m}_2 - \vec{m}_1)$$

where

$$\mathbf{S}_W = \sum_{n \in C_1} (\vec{x}_n - \vec{m}_1)(\vec{x}_n - \vec{m}_1)^T + \sum_{n \in C_2} (\vec{x}_n - \vec{m}_2)(\vec{x}_n - \vec{m}_2)^T$$

- However, while this finds the direction ( $\vec{w}$ ) of decision boundary. Must still solve for  $w_0$  to find the threshold.
- Can be expanded to multiple classes

# Logistic Regression (Discrimination)



# Logistic Regression (Discrimination)

---

- Discriminant model but well-grounded in probability
- Flexible assumptions (exponential family class-conditional densities)
- Differentiable error function (“cross entropy”)
- Works very well when classes are linearly separable



# Logistic Regression (Discrimination)

- Probabilistic discriminative model
- Models posterior probability  $p(C_1|\vec{x})$
- To see this, let's start with the 2-class formulation:

$$\begin{aligned} p(C_1|x) &= \frac{p(\vec{x}|C_1)p(C_1)}{p(\vec{x}|C_1)p(C_1) + p(\vec{x}|C_2)p(C_2)} \\ &= \frac{1}{1 + \exp\left(-\log \frac{p(\vec{x}|C_1)p(C_1)}{p(\vec{x}|C_2)p(C_2)}\right)} \\ &= \frac{1}{1 + \exp(-\alpha)} \quad \text{logistic sigmoid function} \\ &= \sigma(\alpha) \end{aligned}$$


where

$$\alpha = \log \frac{p(\vec{x}|C_1)p(C_1)}{p(\vec{x}|C_2)p(C_2)}$$

# Logistic Regression (Discrimination)

logistic sigmoid function

$$\sigma(\alpha) = \frac{1}{1 + \exp(-\alpha)}$$



“Squashing function” that maps  $(-\infty, +\infty) \rightarrow (0, 1)$

# Logistic Regression (Discrimination)

---

For exponential family of densities,

$$\alpha = \log \frac{p(\vec{x} | C_1)p(C_1)}{p(\vec{x} | C_2)p(C_2)}$$

is a linear function of  $\mathbf{x}$ .

Therefore we can model the posterior probability as a logistic sigmoid acting on a linear function of the attribute vector, and simply solve for the weight vector  $\mathbf{w}$  (e.g. treat it as a discriminant model):

$$y = p(C_1 | \vec{x}) = \sigma(w_0 + \sum_{i=1}^k w_i a_i) \qquad p(C_2 | \vec{x}) = 1 - p(C_1 | \vec{x})$$

$$\text{To classify: } h(\vec{x}_i) = \begin{cases} C_1 & y_i > 0.5 \\ C_2 & o.w. \end{cases}$$